

KUO-HAN HUNG

✉ khhung@stanford.edu 🏠 [khhung-906.github.io](https://github.com/khhung-906) 🎓 [Kuo-Han Hung](#) 🔗 [khhung-906](#)

Education

Stanford University

June 2027 (Expected)

M.S. in Computer Science

Stanford, CA

- GPA: 4.0/4.0; Specialization: Artificial Intelligence and Robotics.

National Taiwan University

Dec. 2024

B.S. in Computer Science and Information Engineering

Taipei, Taiwan

- GPA 4.24/4.3; Participated in Creativity and Entrepreneurship Program.
- Recipient of the Bachelor Thesis Dean's Award, two Presidential Awards, and three Dean's List honors.

Selected Publication

(* indicates equal contribution)

[1] Reinforcement Learning for Real-Time Vision-Language-Action Policies

Perry Dong*, **Kuo-Han Hung***, Dorsa Sadigh, and Chelsea Finn
Under Review

[2] EXPO-FT: Sample-Efficient Reinforcement Learning Finetuning for Vision-Language-Action Models

Perry Dong*, **Kuo-Han Hung***, Tian Gao, Dorsa Sadigh, and Chelsea Finn
Conference on Robot Learning (CoRL), 2026

[3] TQL: Scaling Q-Functions with Transformers by Preventing Attention Collapse

Perry Dong*, **Kuo-Han Hung***, Alexander Swerdlow, Dorsa Sadigh, and Chelsea Finn
International Conference on Machine Learning (ICML), 2026

[4] VICtoR: Learning Hierarchical Vision-Instruction Rewards for Long-horizon Manipulation

Kuo-Han Hung*, Pang-Chi Lo*, Jia-Fong Yeh*, Han-Yuan Hsu, Yi-Ting Chen, and Winston Hsu
International Conference on Learning Representations (ICLR), 2025

[5] Learning Skill Abstraction from Action-Free Videos via Optical Flow

Hung-Chieh Fang*, **Kuo-Han Hung***, Chu-Rong Chen, ..., Min-Hung Chen, and Shao-Hua Sun
ICML Workshop on Building Physically Plausible World Models, 2025

[6] Attention Tracker: Detecting Prompt Injection Attacks in LLMs

Kuo-Han Hung, Ching-Yun Ko, Amrith Rawat, I-Hsin Chung, Winston Hsu and Pin-Yu Chen
Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL), 2025

[7] AED: Adaptable Error Detection for Few-shot Imitation Policy

Jia-Fong Yeh, **Kuo-Han Hung**, Pang-Chi Lo, Chi-Ming Chung, ..., Yi-Ting Chen, and Winston Hsu
Conference on Neural Information Processing Systems (NeurIPS), 2024

Research Experience

Stanford Artificial Intelligence Laboratory (ILIAD & IRIS lab)

Sept. 2025 – Present

Graduate Researcher, Advisor: Prof. Dorsa Sadigh & Prof. Chelsea Finn

Stanford, CA

- Researching **visual representation pretraining** for reinforcement learning.
- Researched **RL finetuning of real-time VLA policies** under inference latency. [1]
 - Decoupled slow VLA action proposals from a fast edit policy that reacts to the latest observation.
 - Raised success from 42% to 97% on four dynamic real-world tasks within 10 min of robot data.
- Researched **sample-efficient RL finetuning of pretrained VLAs** in the real world. [2]

- Built a stable RL finetuning system with residual action edits and Q-guided sampling.
- Reached 30/30 success on all tasks with 19 min of online data on average; code open-sourced.
- Researched **scaling Q-functions** with Transformers in offline RL. [3]
 - Introduced attention entropy regularization to stabilize and scale transformer value functions.
 - Showed prior methods fail to scale, while TQL gains 43% performance with increased capacity.

Embodied AI Lab

June 2025 – Sept. 2025

Undergraduate Researcher, Advisor: Prof. Tsung-Wei Ke

Taipei, Taiwan

- Researched **generalized dexterous manipulation policies** via digital twins.
 - Built a pipeline to convert videos into digital twins using mesh generation and object/human pose estimation.
 - Trained residual RL policies on digital twins from in-the-wild videos, leading to diverse dexterous policies.

Robot Learning Lab

Oct. 2024 – May 2025

Undergraduate Researcher, Advisor: Prof. Shao-Hua Sun

Taipei, Taiwan

- Researched **manipulation policies by learning from videos**. [5]
 - Designed a transformer framework to learn skills from action-free videos using optical flow as action surrogate.
 - Enabled multi-task and long-horizon planning by translating flow-based skill plans into executable actions.

Communication and Multimedia Lab

May. 2022 – May 2024

Undergraduate Researcher, Advisors: Prof. Winston H. Hsu & Prof. Yi-Ting Chen

Taipei, Taiwan

- Researched **vision-instruction correlation rewards** for long-horizon manipulation. [4]
 - Fine-tuned LLMs and VLMs into a hierarchical video reward model for generating stage-aware signals.
 - Achieved a 43% improvement in success rate over state-of-the-art baselines on long-horizon tasks.
- Researched **behavior error detection in few-shot imitation policies** using pattern feature analysis. [7]

Work Experience

IBM Research

June 2024 – Sept. 2024

Research Intern, Host: Dr. Pin-Yu Chen

Yorktown Heights, NY

- Identified the distraction effect of attention scores during **prompt injection attacks on LLM agents**. [6]
- Leveraged mechanistic interpretability to identify attention heads responsible for the distraction effect.
- Invented an attention-based detector, achieving a 31% AUROC improvement over prior best solutions.

Microsoft

July 2023 – June 2024

Applied Scientist Intern, Host: Dr. Chiquan Zhang

Remote, team in Mountain View, CA

- **Integrated Bing Maps with Microsoft Copilot** for conversational map queries with multi-destination search.
- Formulated a pathfinding algorithm that learns from history, **achieving a 10x speedup over Dijkstra**. [??]
- Developed a search system powered by LLMs that adapts to user intent and geographic context.

Teaching Experience

CS 5043: Machine Learning (graduate level)

Teaching Assistant

- Co-designed ML algorithm homework on generalization theory and algorithm understanding (ex. SVM, Random Forest, VC dimension) and a final project about ordinal ranking problems for 250+ students.

Technical Skills

- **Programming Languages:** Python, C/C++, ROS, JavaScript, Golang, SQL, HTML/CSS.
- **ML Knowledge:** Transformers, VLM, LLM, RL, Diffusion, Explainable AI, Regression.
- **ML Toolkits:** PyTorch, JAX, Numpy, Scikit-Learn, Pandas, Matplotlib.
- **Others:** Linux, Docker, Git, Azure, Bash, React, MongoDB.