

# KUO-HAN HUNG

☎ 650-283-6084   ✉ [khhung@stanford.edu](mailto:khhung@stanford.edu)   🏠 [Website](#)   🎓 [Scholar](#)   📞 [khhung-906](#)   🔗 [LinkedIn](#)

## Education

---

### Stanford University

*M.S. in Computer Science*

June 2027 (Expected)

Stanford, CA

- GPA: 4.0/4.0; Specialization: Artificial Intelligence and Robotics.

### National Taiwan University

*B.S. in Computer Science and Information Engineering*

Dec. 2024

Taipei, Taiwan

- GPA: 4.24/4.3, Top 3% of Class. **Additional Special Program:** Creativity and Entrepreneurship (one-year).

## Technical Skills

---

- **Robotics:** Real-Robot Deployment and Training, Real-to-Sim, Real-Time Control, Manipulation, ROS.
- **Parallel Computing:** Distributed Training, Ray, Multiprocessing/Asyncio, Vectorized Environments.
- **Languages & Frameworks:** Python, C/C++, Golang, JavaScript, SQL, Bash, PyTorch, JAX, NumPy.
- **Infrastructure:** Linux, Docker, SLURM, Git, GCP, Azure.

## Publications

(\* indicates equal contribution)

- [1] **Reinforcement Learning for Real-Time Vision-Language-Action Policies**  
Perry Dong\*, **Kuo-Han Hung\***, Dorsa Sadigh, and Chelsea Finn  
Under Review
- [2] **EXPO-FT: Sample-Efficient Reinforcement Learning Finetuning for Vision-Language-Action Models**  
Perry Dong\*, **Kuo-Han Hung\***, Tian Gao, Dorsa Sadigh, and Chelsea Finn  
Conference on Robot Learning (CoRL), 2026
- [3] **TQL: Scaling Q-Functions with Transformers by Preventing Attention Collapse**  
Perry Dong\*, **Kuo-Han Hung\***, Alexander Swerdlow, Dorsa Sadigh, and Chelsea Finn  
International Conference on Machine Learning (ICML), 2026
- [4] **VICtoR: Learning Hierarchical Vision-Instruction Rewards for Long-horizon Manipulation**  
**Kuo-Han Hung\***, Pang-Chi Lo\*, Jia-Fong Yeh\*, Han-Yuan Hsu, Yi-Ting Chen, and Winston Hsu  
International Conference on Learning Representations (ICLR), 2025
- [5] **Learning Skill Abstraction from Action-Free Videos via Optical Flow**  
Hung-Chieh Fang\*, **Kuo-Han Hung\***, Chu-Rong Chen, ..., Min-Hung Chen, and Shao-Hua Sun  
ICML Workshop on Building Physically Plausible World Models, 2025
- [6] **DexMan: Learning Bimanual Dexterous Manipulation from Human and Generated Videos**  
Jhen Hsieh, Kuan-Hsun Tu, **Kuo-Han Hung**, and Tsung-Wei Ke  
Under Review
- [7] **Attention Tracker: Detecting Prompt Injection Attacks in LLMs**  
**Kuo-Han Hung**, Ching-Yun Ko, Ambrish Rawat, I-Hsin Chung, Winston Hsu, and Pin-Yu Chen  
Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL), 2025
- [8] **AED: Adaptable Error Detection for Few-shot Imitation Policy**  
Jia-Fong Yeh, **Kuo-Han Hung**, Pang-Chi Lo, Chi-Ming Chung, ..., Yi-Ting Chen, and Winston Hsu  
Conference on Neural Information Processing Systems (NeurIPS), 2024

## Research Experience

---

### Stanford Artificial Intelligence Laboratory (ILIAD & IRIS Labs)

*Graduate Researcher, Advisors: Prof. Dorsa Sadigh & Prof. Chelsea Finn*

Sept. 2025 – Present

Stanford, CA

- **RL finetuning of real-time VLA policies** under inference latency. [1]
  - Decoupled slow VLA proposals from a fast edit policy reacting to the latest observation, raising success from 42% to 97% across six dexterous tasks including light-string routing and flower insertion, with under 10 min of online data.
  - Built the full real-time policy stack enabling effective RL inference and training: background-thread VLA inference overlapped with chunk execution, action-prefix inpainting, and a 30 Hz Cartesian control loop.
- **Sample-efficient RL finetuning of pretrained VLAs** in the real world. [2]
  - Stabilized RL finetuning with residual action edits and Q-guided sampling, reaching 30/30 success across six dexterous real-world tasks including light-string routing and flower insertion, with 19 min of online data.
  - Designed and built the real-world RL pipeline: asynchronous multi-GPU training, WebSocket server-client rollout, replay buffer, data pipelines, and DROID integration (100+ GitHub stars).
- **Scaling Q-functions** with Transformers in offline RL. [3]
  - Introduced attention entropy regularization to stabilize and scale Transformer value function training.
  - Showed that all prior methods fail to scale, while TQL gains 43% success rate as model size increases.
- **Visual representation pretraining** for reinforcement learning (ongoing).

## Embodied AI Lab

June 2025 – Sept. 2025

Undergraduate Researcher, Advisor: Prof. Tsung-Wei Ke

Taipei, Taiwan

- **Generalized dexterous manipulation policies** via digital twins. [6]
  - Built a real-to-sim pipeline: videos to digital twins via mesh generation and object/human pose estimation.
  - Trained residual RL policies on the twins, raising success rate 19% over prior methods on OakInk-v2.

## Robot Learning Lab

Oct. 2024 – May 2025

Undergraduate Researcher, Advisor: Prof. Shao-Hua Sun

Taipei, Taiwan

- **Manipulation policies learned from action-free videos.** [5]
  - Designed a Transformer to learn skills from action-free videos, using optical flow as an action surrogate.
  - Enabled multi-task and long-horizon planning, raising success rate 21% over the best MetaWorld baseline.

## Communication and Multimedia Lab

May 2022 – May 2024

Undergraduate Researcher, Advisors: Prof. Winston H. Hsu & Prof. Yi-Ting Chen

Taipei, Taiwan

- **Vision-instruction correlation rewards** for long-horizon manipulation. [4]
  - Finetuned LLMs and VLMs into a hierarchical video reward model for generating stage-aware signals.
  - Achieved a 43% improvement in success rate over state-of-the-art baselines on long-horizon tasks.
- **Behavior error detection** in few-shot imitation policies via representation pattern analysis. [8]

## Work Experience

### IBM Research

June 2024 – Sept. 2024

Research Intern, Host: Dr. Pin-Yu Chen

Yorktown Heights, NY

- Found that attention drifts from the original instruction during **prompt injection attacks on LLM agents.** [7]
- Traced the effect to specific attention heads using mechanistic interpretability.
- Designed an attention-based detector, a 31% AUROC improvement over the prior state of the art.

### Microsoft

July 2023 – June 2024

Applied Scientist Intern, Host: Dr. Chiqun Zhang

Mountain View, CA (Remote)

- **Integrated Bing Maps with Microsoft Copilot** for conversational, multi-destination map queries.
- Formulated a pathfinding algorithm that learns from history, **achieving a 10× speedup over Dijkstra.**
- Developed an LLM-powered search system that adapts to user intent and geographic context.

## Awards & Honors

- **Student Research Scholarship**, Taiwan National Science Council. Sept. 2024
- **Bachelor Thesis Dean's Award**, National Taiwan University — top 2 theses in the CS department. May 2024
- **Lin Hsiung Chen Memorial Scholarship** — for top 1 CS student in NTU. Nov. 2023
- **Creativity and Entrepreneurship Excellence Award**, NTU D-School. June 2023
- **Presidential Award** (×2), National Taiwan University — top 2% of the CS department. Sept. 2022, 2023
- **Dean's List** (×3), National Taiwan University — top 5% of the CS department. June 2021, 2022, 2023